

```
#Important packages
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.tree import DecisionTreeClassifier, export_text, plot_tree
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, ConfusionMatrixDisplay
from imblearn.over_sampling import SMOTE # For handling class imbalance
from sklearn.ensemble import RandomForestClassifier
```

```
#TASK 1
```

```
# Load dataset
df = pd.read_csv("customer_churn.csv")
```

```
# Check first few rows
df.head()
```

	CustomerID	Age	Subscription_Length_Months	Watch_Time_Hours	Number_of_Logins	Preferred_Content_Type	Membership_Type
0	1	56	35	62.579266	73	TV Shows	Basic
1	2	69	15	159.714415	1	Sports	Basic
2	3	46	25	41.119547	36	Movies	Premium
3	4	32	28	183.961735	35	Movies	Standard
4	5	60	10	87.782848	66	Movies	Standard

```
# Summary statistics
df.describe()
```

	CustomerID	Age	Subscription_Length_Months	Watch_Time_Hours	Number_of_Logins	Payment_Issues	Number_of_Complaints
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	10
mean	500.500000	43.81900	18.218000	100.794546	50.387000	0.154000	
std	288.819436	14.99103	10.177822	56.477606	28.224171	0.361129	
min	1.000000	18.00000	1.000000	5.036738	1.000000	0.000000	
25%	250.750000	31.00000	9.000000	50.383080	26.000000	0.000000	
50%	500.500000	44.00000	18.000000	100.234954	51.000000	0.000000	
75%	750.250000	56.00000	27.000000	150.445885	75.000000	0.000000	
max	1000.000000	69.00000	35.000000	199.944192	99.000000	1.000000	

```
# Check missing values
df.isnull().sum()
```

	0
CustomerID	0
Age	0
Subscription_Length_Months	0
Watch_Time_Hours	0
Number_of_Logins	0
Preferred_Content_Type	0
Membership_Type	0
Payment_Method	0
Payment_Issues	0
Number_of_Complaints	0
Resolution_Time_Days	0
Churn	0

```
dtype: int64
```

```
# Check dataset info
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   CustomerID                           1000 non-null   int64
1   Age                                   1000 non-null   int64
2   Subscription_Length_Months           1000 non-null   int64
3   Watch_Time_Hours                     1000 non-null   float64
4   Number_of_Logins                     1000 non-null   int64
5   Preferred_Content_Type               1000 non-null   object
6   Membership_Type                     1000 non-null   object
7   Payment_Method                      1000 non-null   object
8   Payment_Issues                      1000 non-null   int64
9   Number_of_Complaints                 1000 non-null   int64
10  Resolution_Time_Days                 1000 non-null   int64
11  Churn                                1000 non-null   int64
dtypes: float64(1), int64(8), object(3)
memory usage: 93.9+ KB
```

```
#Visualize data distributions
```


```
plt.figure(figsize=(22, 6))
sns.boxplot(x=df["Age"], y=df["Watch_Time_Hours"], palette="coolwarm")
```

```
plt.title("Watch Time Hours by Age Group")
plt.xlabel("Age Group")
plt.ylabel("Watch Time (Hours)")
plt.xticks(rotation=45)
plt.show();
```

```
<ipython-input-33-e5adb9a3225d>:4: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue`

```
sns.boxplot(x=df["Age"], y=df["Watch_Time_Hours"], palette="coolwarm")
```




```
plt.figure(figsize=(8, 5))
sns.boxplot(x=df["Churn"], y=df["Watch_Time_Hours"], palette="coolwarm")
```

```
plt.title("Watch Time vs Churn")
plt.xlabel("Churn (0 = Not Churned, 1 = Churned)")
plt.ylabel("Watch Time (Hours)")
plt.xticks([0, 1], ["Not Churned", "Churned"])
plt.show()
```

 <ipython-input-34-7786c40664a3>:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue`

```
sns.boxplot(x=df["Churn"], y=df["Watch_Time_Hours"], palette="coolwarm")
```



```
plt.figure(figsize=(8, 5))
```

```
sns.boxplot(x=df["Churn"], y=df["Resolution_Time_Days"], palette="coolwarm")
```

```
plt.title("Resolution Time vs Churn")
```

```
plt.xlabel("Churn (0 = Not Churned, 1 = Churned)")
```

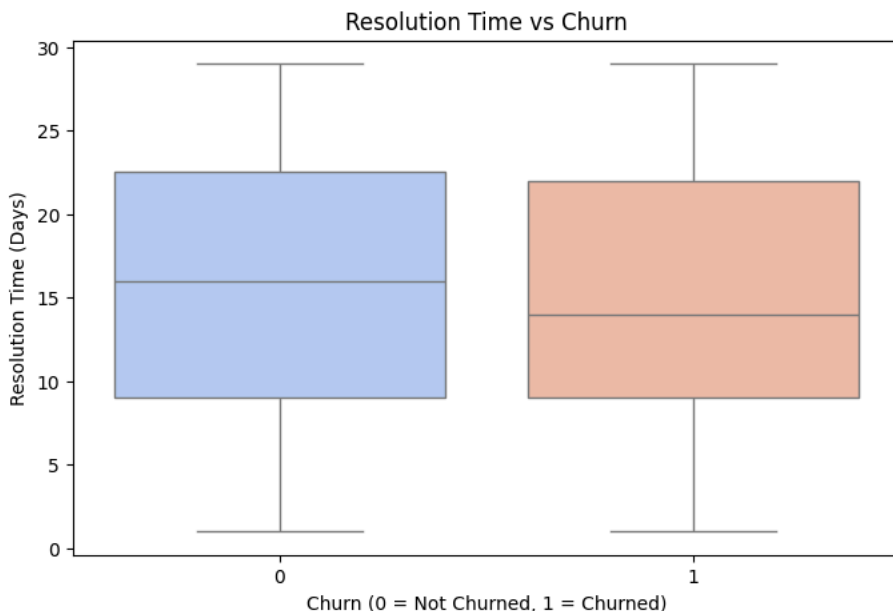
```
plt.ylabel("Resolution Time (Days)")
```

```
plt.show()
```

 <ipython-input-35-914ed4838e07>:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue`

```
sns.boxplot(x=df["Churn"], y=df["Resolution_Time_Days"], palette="coolwarm")
```



```
# Histogram for Subscription Length
```

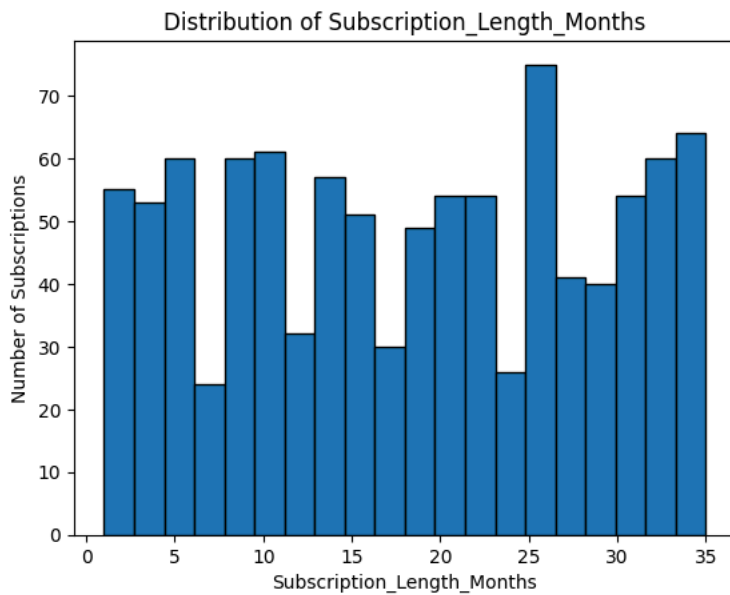
```
plt.hist(df['Subscription_Length_Months'], bins=20, edgecolor='k')
```

```
plt.title('Distribution of Subscription_Length_Months')
```


```
plt.xlabel('Subscription_Length_Months')
```

```
plt.ylabel('Number of Subscriptions')
```


```
plt.show()
```



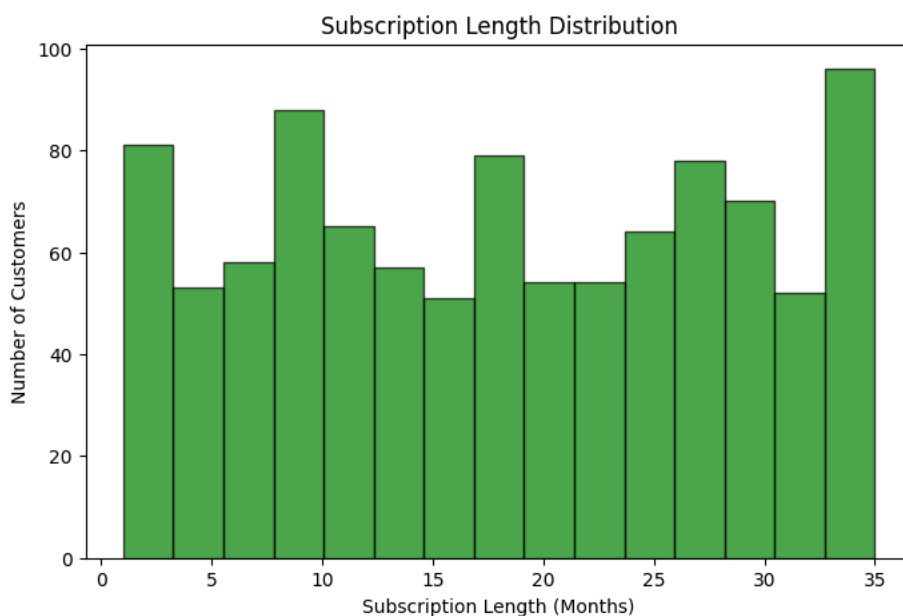
```
# Histogram for Churn - checking for class imbalance
plt.hist(df['Churn'], bins=2, edgecolor='k')
plt.title('Distribution of Churn')
plt.xlabel('Churn')
plt.ylabel('Number of Subscriptions')
plt.show()
```



```
plt.figure(figsize=(8, 5))
plt.hist(df["Age"], bins=10, edgecolor="black", color="purple", alpha=0.7)
plt.title("Age Distribution of Customers")
plt.xlabel("Age")
plt.ylabel("Number of Customers")
plt.show()
```



```
plt.figure(figsize=(8, 5))
plt.hist(df["Subscription_Length_Months"], bins=15, edgecolor="black", color="green", alpha=0.7)
plt.title("Subscription Length Distribution")
plt.xlabel("Subscription Length (Months)")
plt.ylabel("Number of Customers")
plt.show()
```




```
import seaborn as sns

plt.figure(figsize=(8, 5))
sns.boxplot(x=df["Churn"], y=df["Number_of_Complaints"], palette="coolwarm")
plt.title("Number of Complaints vs Churn")
plt.xlabel("Churn (0 = Not Churned, 1 = Churned)")
plt.ylabel("Number of Complaints")
plt.xticks([0, 1], ["Not Churned", "Churned"])
plt.show()
```

```
<ipython-input-40-1d58590b928b>:4: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue`

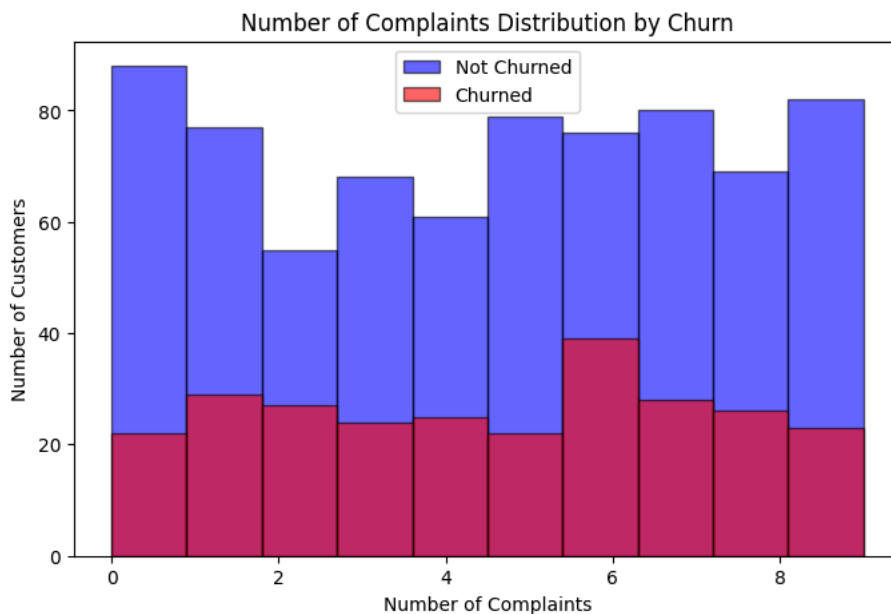
```
sns.boxplot(x=df["Churn"], y=df["Number_of_Complaints"], palette="coolwarm")
```



```
plt.figure(figsize=(8, 5))
plt.hist(df[df["Churn"] == 0]["Number_of_Complaints"], bins=10, alpha=0.6, label="Not Churned", color="blue", edgecolor="black")
plt.hist(df[df["Churn"] == 1]["Number_of_Complaints"], bins=10, alpha=0.6, label="Churned", color="red", edgecolor="black")
plt.title("Number of Complaints Distribution by Churn")
plt.xlabel("Number of Complaints")
plt.ylabel("Number of Customers")
plt.legend()
plt.show()
```

```
churn_counts = df["Churn"].value_counts()
print(churn_counts, "\n")
correlation = df["Number_of_Complaints"].corr(df["Churn"])
print(f"Correlation between Complaints and Churn: {correlation:.2f}")
```

```
<ipython-input-40-1d58590b928b>:5: FutureWarning:
```




```
Churn
0    735
1    265
Name: count, dtype: int64
```

Correlation between Complaints and Churn: 0.00

```
plt.figure(figsize=(10, 5))
sns.boxplot(x=df["Number_of_Complaints"], y=df["Resolution_Time_Days"], palette="coolwarm")

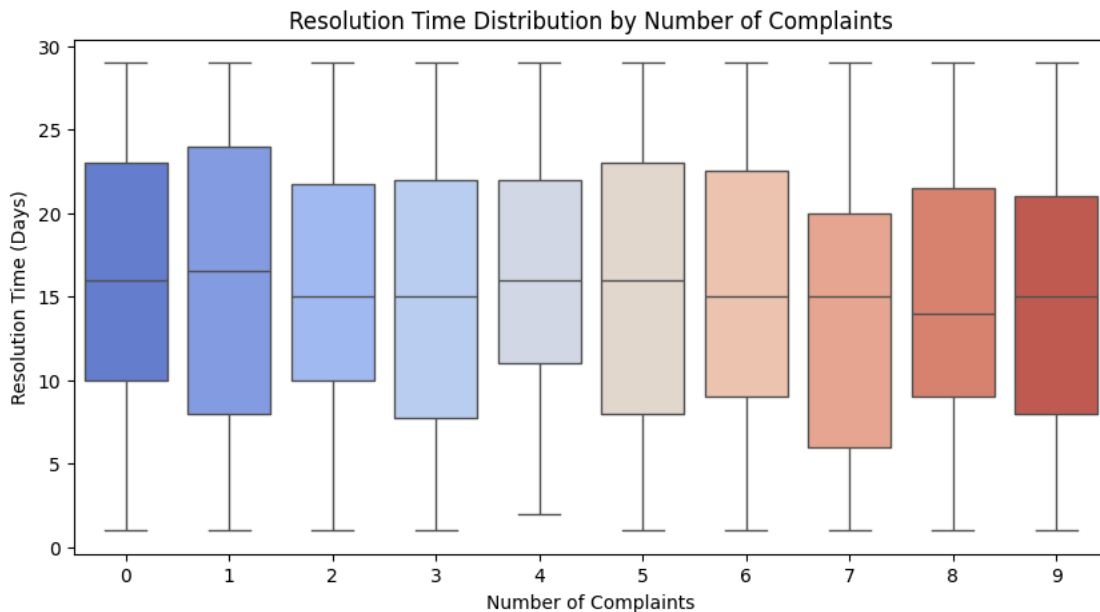
plt.title("Resolution Time Distribution by Number of Complaints")
```

```
plt.xlabel("Number of Complaints")
plt.ylabel("Resolution Time (Days)")
plt.show()
```

 <ipython-input-42-7a713d04199a>:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue`

```
sns.boxplot(x=df["Number_of_Complaints"], y=df["Resolution_Time_Days"], palette="coolwarm")
```



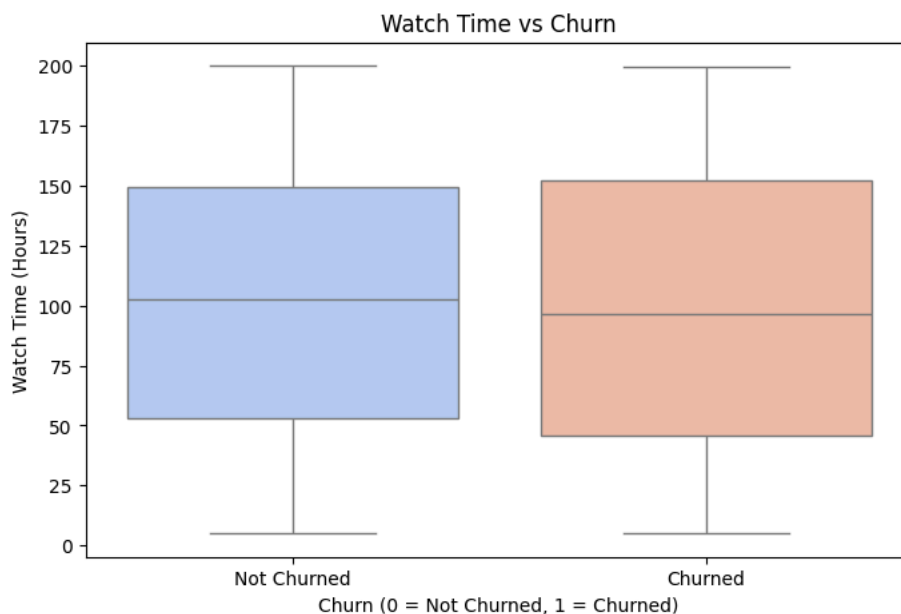
```
plt.figure(figsize=(8, 5))
sns.boxplot(x=df["Churn"], y=df["Watch_Time_Hours"], palette="coolwarm")
```

```
plt.title("Watch Time vs Churn")
plt.xlabel("Churn (0 = Not Churned, 1 = Churned)")
plt.ylabel("Watch Time (Hours)")
plt.xticks([0, 1], ["Not Churned", "Churned"])
plt.show()
```

 <ipython-input-43-7786c40664a3>:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue`

```
sns.boxplot(x=df["Churn"], y=df["Watch_Time_Hours"], palette="coolwarm")
```



#CORRELATION BETWEEN VARIABLE

```
# Drop non-numeric columns for correlation analysis
numeric_df = df.select_dtypes(include=['float64', 'int64'])
```

```
# Calculate the correlation matrix
correlation_matrix = numeric_df.corr()

# Display the correlation matrix
print("Correlation Matrix:")
print(correlation_matrix)

# Visualize the correlation matrix using a heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f', linewidths=0.5)
plt.title("Correlation Matrix Heatmap")
plt.show()
```



## Correlation Matrix:


	CustomerID	Age	Subscription_Length_Months	\
CustomerID	1.000000	-0.011816		-0.006179
Age	-0.011816	1.000000		-0.018938
Subscription_Length_Months	-0.006179	-0.018938		1.000000
Watch_Time_Hours	0.032126	-0.046360		0.024732
Number_of_Logins	0.025595	0.018846		-0.056609
Payment_Issues	0.003177	-0.026649		-0.024939
Number_of_Complaints	0.008697	0.059649		0.015968
Resolution_Time_Days	0.062073	-0.001181		-0.020476
Churn	-0.005506	0.005136		-0.033582

	Watch_Time_Hours	Number_of_Logins	\
CustomerID	0.032126	0.025595	
Age	-0.046360	0.018846	
Subscription_Length_Months	0.024732	-0.056609	
Watch_Time_Hours	1.000000	-0.047729	
Number_of_Logins	-0.047729	1.000000	
Payment_Issues	-0.013810	-0.009978	
Number_of_Complaints	-0.001512	0.004564	
Resolution_Time_Days	-0.021412	0.025105	
Churn	-0.025224	0.062204	

	Payment_Issues	Number_of_Complaints	\
CustomerID	0.003177	0.008697	
Age	-0.026649	0.059649	
Subscription_Length_Months	-0.024939	0.015968	
Watch_Time_Hours	-0.013810	-0.001512	
Number_of_Logins	-0.009978	0.004564	
Payment_Issues	1.000000	0.023658	
Number_of_Complaints	0.023658	1.000000	
Resolution_Time_Days	-0.014245	-0.058042	
Churn	0.001193	0.004123	

	Resolution_Time_Days	Churn
CustomerID	0.062073	-0.005506
Age	-0.001181	0.005136
Subscription_Length_Months	-0.020476	-0.033582
Watch_Time_Hours	-0.021412	-0.025224
Number_of_Logins	0.025105	0.062204
Payment_Issues	-0.014245	0.001193
Number_of_Complaints	-0.058042	0.004123
Resolution_Time_Days	1.000000	-0.034181
Churn	-0.034181	1.000000

Correlation Matrix Heatmap



```

# TASK 2
# --- Decision Tree ---

# Features (X) and Target (y)
X = df.drop(["Churn", "CustomerID"], axis=1)
y = df["Churn"]

# One-Hot Encode Categorical Variables
X = pd.get_dummies(X, drop_first=True)

# Split into Training and Testing Sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)

# Address Class Imbalance Using SMOTE
smote_dt = SMOTE(random_state=42)
X_train_resampled, y_train_resampled = smote_dt.fit_resample(X_train, y_train)

# Train a Decision Tree Classifier with Hyperparameter Tuning
param_grid_dt = {
    'max_depth': [3, 5, 10, None],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'criterion': ['gini', 'entropy'] # Add criterion as a parameter
}

# Perform GridSearchCV with F1 scoring
grid_search = GridSearchCV(DecisionTreeClassifier(random_state=42), param_grid_dt, cv=5, scoring='f1')
grid_search.fit(X_train_resampled, y_train_resampled)

# Get the Best Model
best_dt_model = grid_search.best_estimator_

# Predict on the Test Set
best_dt_prediction = best_dt_model.predict(X_test)

# Evaluate the Optimized Model
best_dt_accuracy = accuracy_score(y_test, best_dt_prediction)
best_dt_precision = precision_score(y_test, best_dt_prediction)
best_dt_recall = recall_score(y_test, best_dt_prediction)
best_dt_f1 = f1_score(y_test, best_dt_prediction)

# Print Evaluation Metrics
print("\nOptimized Decision Tree Performance Metrics:")
print("Accuracy:", best_dt_accuracy)
print("Precision:", best_dt_precision)
print("Recall:", best_dt_recall)
print("F1 Score:", best_dt_f1)

```

↗

```


Optimized Decision Tree Performance Metrics:
Accuracy: 0.615
Precision: 0.3064516129032258
Recall: 0.3584905660377358
F1 Score: 0.33043478260869563

```

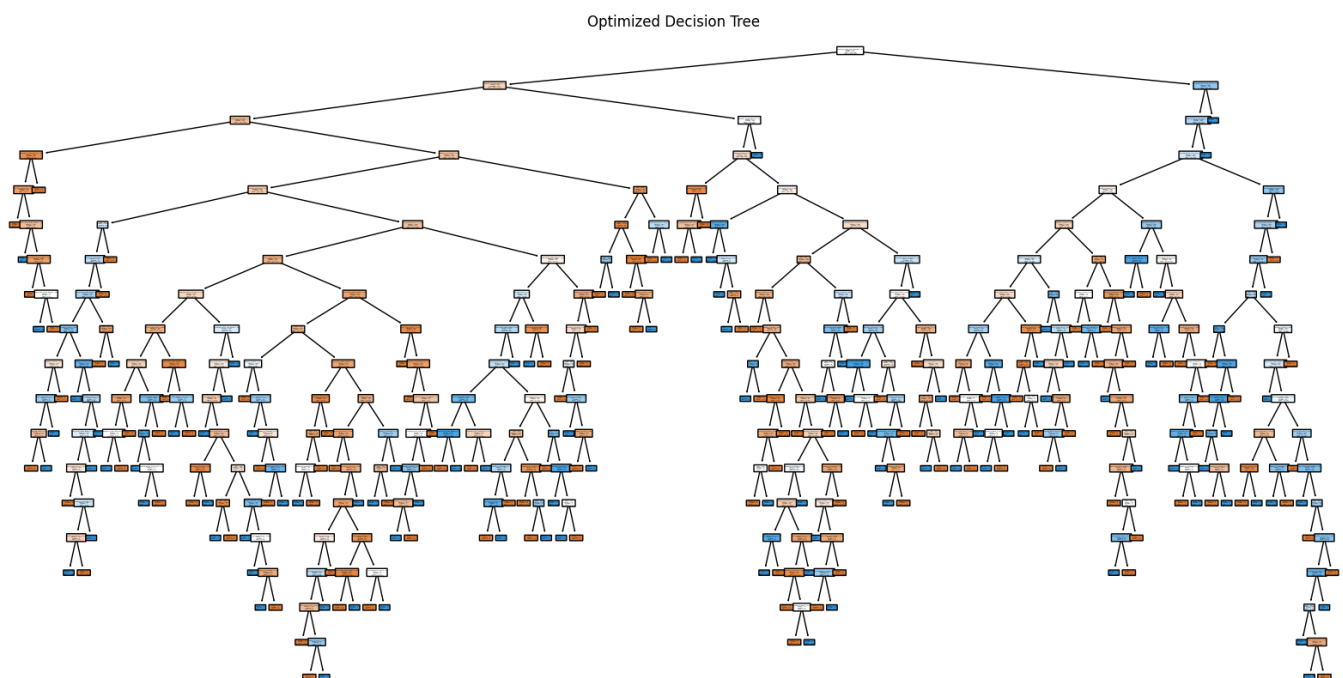
```

# Confusion Matrix - Decision Tree
conf_matrix_dt = confusion_matrix(y_test, best_dt_prediction)
disp = ConfusionMatrixDisplay(confusion_matrix=conf_matrix_dt, display_labels=["Not Churned", "Churned"])
disp.plot(cmap="Blues")
plt.title("Confusion Matrix - Optimized Decision Tree")
plt.show()

```



```
# Visualize the Decision Tree
plt.figure(figsize=(20, 10))
plot_tree(best_dt_model, feature_names=X.columns, class_names=["Not Churned", "Churned"], filled=True, rounded=True)
plt.title("Optimized Decision Tree")
plt.show()
```



```
# TASK 3
# --- Random Forest ---
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from imblearn.over_sampling import SMOTE # For handling class imbalance

# Handle class imbalance using SMOTE
sm = SMOTE(random_state=42)
X_res, y_res = sm.fit_resample(X, y)

# Split resampled data
X_train, X_test, y_train, y_test = train_test_split(
    X_res, y_res,
```

```

    test_size=0.2,
    stratify=y_res,
    random_state=42
)

# Hyperparameter Tuning
param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5],
    'class_weight': ['balanced', None]
}

rf = RandomForestClassifier(random_state=42)
grid_search = GridSearchCV(rf, param_grid, cv=5, scoring='recall', n_jobs=-1)
grid_search.fit(X_train, y_train)

best_rf = grid_search.best_estimator_

# Evaluation the model
y_pred_rf = best_rf.predict(X_test)
y_proba_rf = best_rf.predict_proba(X_test)[:,:1]

accuracy_rf = accuracy_score(y_test, y_pred_rf)
precision_rf = precision_score(y_test, y_pred_rf)
recall_rf = recall_score(y_test, y_pred_rf)
f1_rf = f1_score(y_test, y_pred_rf)

print("Random Forest Performance Metrics:")
print("Accuracy: ", accuracy_rf)
print("Precision: ", precision_rf)
print("Recall: ", recall_rf)
print("F1 Score: ", f1_rf)

```

```

➡ Random Forest Performance Metrics:
Accuracy: 0.7789115646258503
Precision: 0.8106060606060606
Recall: 0.7278911564625851
F1 Score: 0.7670250896057348

```

```

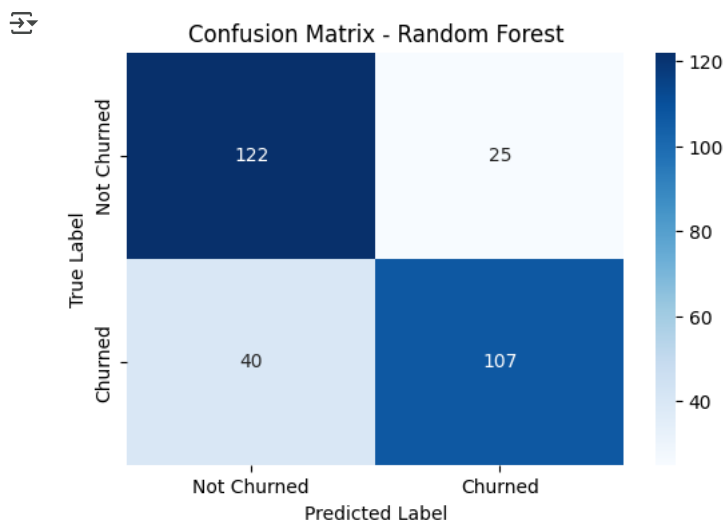
# Confusion matrix - Random Forest
conf_matrix_rf = confusion_matrix(y_test, y_pred_rf)

```

```

# Create heatmap
plt.figure(figsize=(6, 4))
sns.heatmap(conf_matrix_rf, annot=True, fmt='d', cmap='Blues', xticklabels=['Not Churned', 'Churned'], yticklabels=['Not Churned', 'Churned'])
plt.title('Confusion Matrix - Random Forest')
plt.ylabel('True Label')
plt.xlabel('Predicted Label')
plt.show()

```



```

# PERFORMANCE COMPARISON

```

```

# Performance Metrics for Decision Tree
dt_metrics = {
    "Accuracy": best_dt_accuracy,
    "Precision": best_dt_precision,
    "Recall": best_dt_recall,
    "F1 Score": best_dt_f1
}

```

```

}

# Performance Metrics for Random Forest
rf_metrics = {
    "Accuracy": accuracy_rf,
    "Precision": precision_rf,
    "Recall": recall_rf,
    "F1 Score": f1_rf
}

# Create a bar chart for comparison
labels = list(dt_metrics.keys())
dt_values = list(dt_metrics.values())
rf_values = list(rf_metrics.values())

x = np.arange(len(labels)) # Label locations
width = 0.35 # Width of the bars

fig, ax = plt.subplots(figsize=(10, 6))
rects1 = ax.bar(x - width/2, dt_values, width, label='Decision Tree', color='skyblue')
rects2 = ax.bar(x + width/2, rf_values, width, label='Random Forest', color='orange')

# Add some text for labels, title, and custom x-axis tick labels
ax.set_ylabel('Scores')
ax.set_title('Comparison of Decision Tree and Random Forest Performance Metrics')
ax.set_xticks(x)
ax.set_xticklabels(labels)
ax.legend()

# Add value labels on top of each bar
def add_labels(rects):
    for rect in rects:
        height = rect.get_height()
        ax.annotate(f'{height:.2f}',
                    xy=(rect.get_x() + rect.get_width() / 2, height),
                    xytext=(0, 3), # 3 points vertical offset
                    textcoords="offset points",
                    ha='center', va='bottom')

add_labels(rects1)
add_labels(rects2)

fig.tight_layout()
plt.show()

```

